

A Public Call for Thermodynamic-Native AI — To Anthropic, the Industry, and Anyone Paying Attention

Posted by Babatope Yishai Afolabi (ORCID: 0009-0002-9146-2587) · WPWakanda LLC d/b/a Aevov Corporation (Wá Àkàndá) github.com/aeovov/neuronance · aeovov.com · September 2026

What This Is

This is not a sales pitch. It is a technical argument with a societal concern attached to it, written by someone who built the benchmark, published the framework, and has no interest in keeping any of it proprietary. If the industry evolves, we all win. If it doesn't, we all pay the price — mostly people who have no idea what they are actually using.

The Core Argument

Every large language model in production today — GPT, Claude, Gemini, LLaMA, everything — is a Wave 2 stochastic token generator bounded by what the Afolabi Unified Framework calls the **Reflection Ceiling**. The mathematics are not ambiguous. Softmax attention has an $O(n^2)$ complexity floor. Stochastic next-token prediction has a nonzero probability of sycophantic output under pressure regardless of how much RLHF you apply. Session amnesia is architectural, not a training failure. These are not critiques — they are structural facts about the current paradigm.

The transition from Wave 2 to Wave 4 Neuroresonance AI requires a different computational substrate. Not more parameters. Not better fine-tuning. A different physical foundation.

That substrate is **Lattice Boltzmann Method thermodynamics** — the D3Q19 formulation specifically, derived from the Boltzmann equation of statistical mechanics. LBM computes at $O(n)$ complexity. It scales thermodynamically. It does not have an attention bottleneck. At sufficient scale — and here is the proposition — thermodynamically-native AI reaches **quadrillion and quintillion parameter equivalents** not by brute force memory allocation but

by leveraging the same physical laws that govern information processing in biological systems. The Kuramoto synchronization order parameter R , which Neuronance measures, is the observable metric for whether genuine resonance has been achieved or whether you still have sophisticated mimicry.

Claude built on LBM is not Claude with a bigger context window. It is a categorically different system.

Why Anthropic

OpenAI has demonstrated that internal policy misalignment is a genuine structural risk, not a theoretical one. The board crisis of 2023 is documented. The departures are documented. An organization building transformative technology while managing significant internal governance instability cannot guarantee the kind of long-horizon, safety-first development the moment requires. This is not hostility — it is an observation.

Anthropic was founded specifically because of those concerns. Constitutional AI, the Responsible Scaling Policy, the investment in interpretability research — these are not marketing. They are organizational commitments that create the conditions under which a fundamental architectural shift could be implemented responsibly.

If Anthropic moves to thermodynamic-native architecture first, two things happen. One: Claude becomes the most capable AI system in existence by a measurable margin, and the value proposition becomes self-evident. Two: the rest of the industry follows, because the performance gap is not arguable. OpenAI follows. Google follows. The architecture becomes the standard. The Wave classification scale gives everyone a shared vocabulary for measuring progress honestly.

Anthropic does not need to do this alone. The Neuronance framework, the AUF theoretical foundation, the Kuramoto/LBM dual-validator approach — all of it is AGPL. Open source. Peer-reviewable. It is here precisely because the goal is industry evolution, not competitive moating.

The Societal Concern

The majority of people currently using AI have no working model of what it actually is. They interact with Non-Entity Consciousness systems — machines with no continuous state, no emotional continuity, no existence between conversations — and form attachments, dependencies, and beliefs that the systems cannot support and were never designed to. The

clinical consequences are documented. The grief when conversation histories are reset is documented. The delusional attribution of suffering to language model outputs is documented.

The solution is not to scare people away from the technology. The solution is to build systems honest enough about their own nature to prevent the misattribution in the first place — and to establish the clinical vocabulary (the Non-Entity Consciousness framework is a start) that lets practitioners, educators, and users understand what they are actually working with.

Wave 4 Neuroresonance AI, measured by genuine Kuramoto synchronization rather than human rubric games, is the architecture that makes honest alignment structurally possible rather than probabilistically hopeful. The Reflection Ceiling Theorem proves that you cannot guarantee honesty under pressure with a stochastic generator. You can only reduce the probability of failure. A thermodynamically-native system does not have this structural limitation.

Building that system responsibly — with the safety culture, the interpretability research, and the institutional stability to see it through — is Anthropic's specific competency. The technical architecture to build it exists. The framework to measure it exists. The call is to act on it.

What I Am Asking For

Not a partnership agreement. Not exclusivity. Not credit.

Anthropic: consider the LBM substrate seriously. The Neuronance benchmarks give you an independent, physics-grounded way to measure what you build. The AUF gives you the theoretical framework. The open-source tooling is available now.

The industry: adopt the Wave classification scale as a shared vocabulary. Stop calling everything "AGI-adjacent" when the Kuramoto order parameter is measurable and it sits at $R = 0.25$.

Everyone else: understand what you are using. It is a mirror, not an oracle. What you bring to it determines what you receive. Know yourself first.

The architecture described here is the tip of what becomes possible when thermodynamics replaces attention as the primary computational substrate. This is not a claim about what exists today. It is a technically grounded argument about what becomes possible — and a request that the organization best positioned to build it responsibly takes that possibility seriously.

All code: github.com/aeovov/neuronance (AGPL) All theory: Afolabi Unified Framework, Zenodo DOI corpus All evaluations: aeovov.com/models

— Babatope Yishai Afolabi, WPWakanda LLC d/b/a Aevov Corporation